

report_export_case

report_export_case: HIGH kockázat, UPSIZE_REQUIRED procurement band · Többnyire kalibrált
Nagyobb hardver vagy több GPU szükséges a célterheléshez.

Safe 24/7 capacity	1.018 req/s
Latency p95	1549 ms
GPU util p95	83.4%
Gap	-1.982 req/s

Workload összefoglaló

Workload ID	bp_template_rag
Class	retrieval_augmented
Family	enterprise_rag
Arrival rate	3.000 req/s
Planning profile	baseline

Hardware ajánlás

Hardware	NVIDIA DGX H200 — 8× H200
GPU count	8
Runtime	vllm / vllm_continuous_batching
Software stack	vLLM
Deployment	container_on_bare_metal

24/7 safe capacity

Configured arrival	3.000 req/s
Steady-state capacity	1.708 req/s
Safe 24/7 capacity	1.018 req/s
Headroom ratio	0.17
Capacity gap	-1.982 req/s

Bottleneck

Primary bottleneck stage: **generate_answer**

generate_answer	100.0%
-----------------	--------

Baseline vs override vs safe profile

Scenario	Planning	Safe cap	Latency p95	GPU util p95	Band
report_export_case	baseline	1.018	1549 ms	83.4%	UPSIZE_REQUIRED

Assumptions & limits

- Runtime: vllm/vllm_continuous_batching
- Deployment: container_on_bare_metal
- Hardware: nvidia_dgx_h200_8gpu (8 GPU worker domain)
- Most frequent bottleneck stage: generate_answer
- Role bindings: embedding_role=intfloat/multilingual-e5-large, llm_role=mistralai/Mistral-Small-3.1-24B-Instruct-2503, parse_role=nvidia/nv-ingest-doc-splitter, reranker_role=BAAI/bge-reranker-v2-m3, vlm_role=OpenGVLab/InternVL2_5-8B
- Calibration coverage: 3/5 stages matched to benchmark anchors.
- Model compatibility: warning (ok=2, warning=3, error=0).
- Procurement band: UPSIZE_REQUIRED. Nagyobb hardver vagy több GPU szükséges a célterheléshez.
- Configured arrival rate: 3.00/s, safe 24/7 capacity: 1.02/s.
- A pipeline több szakasza benchmark anchorhoz illeszkedett, de nem teljes a lefedettség.